

# TattvaDrishti

## Analyst-Centric Intelligence for **AI-Enabled Malign Information Operations**

Explainable, sovereign decision-support for detecting, correlating and securely sharing intelligence on AI-driven malign information operations.



Team  
**ASHTOJ**



Institute  
Keshav Memorial Institute of  
Technology, Hyderabad



**SIH 2025 Winners** • AICTE Cyber Security  
Internship 2025–2026



# AI is changing information warfare



## Scale of deception

LLMs generate persuasive disinformation, phishing and synthetic narratives at low cost and high volume.



## Coordination & evasion

Campaigns adapt across channels and are difficult to trace with manual or single-signal tools.



## High-stakes consequences


Public trust, elections, crisis response and national security are directly affected.



**Why now:** rapid LLM progress and multilingual, multi-platform threats demand agency-ready defense.

# What TattvaDrishti does

One operational flow from raw content to **actionable intelligence**.



Ingests messages, archives, public-channel content and images



Fuses local reasoning with AI-origin, behavioral and stylometric signals



Surfaces explainable risk and secure intelligence-sharing outputs

# Core capabilities



## 1. Manual & batch intake

Single narratives or large JSON, JSONL and NDJSON archives.



## 2. Telegram & Reddit normalization

Public-channel content is converted into a common analysis schema.



## 3. Explainable multi-signal scoring

Semantic risk, AI-origin, behavioral cues and stylometry.



## 4. Threat graph & actor analysis

Links content, actors, narratives, regions and URLs.



## 5. Keyword triage & geographic risk

Filter investigations and visualize where suspicious narratives appear.



## 6. Secure sharing & audit trail


Redacted, signed, encrypted case packages with traceable lineage.



Built as an **analyst-centric** platform, not a single-score detector.

# Evidence fusion, not single-score detection

Independent evidence is preserved and fused; no single detector becomes the final verdict.



**Human-in-the-loop:** the platform prioritizes review; the final judgement stays with the analyst.

# Built to scale through batch analysis

From one message to large evidence archives.



**5,000**

default record limit



**25 MiB**

upload limit



**16**

default transformer  
batch size



**2**

default concurrent  
Ollama calls



**LATEST INTERNAL DEMONSTRATION:**


**~150 messages in approximately 10–15 seconds**

Team-reported engineering observation on a specific test setup. Not yet a formal benchmark; results vary with hardware, models, message length, warm-up state and local-LLM usage.

# From suspicious messages to coordinated campaigns

A harmful information operation is rarely visible in one sentence. TattvaDrishti treats graph intelligence as a core investigation layer.

-  Actor rankings and summaries
-  Repeated narrative clusters
-  Propagation chains across platforms
-  Keyword and event filtering
-  Geographic and SIEM/CTI correlation



**Guardrail:** coordination does not automatically imply malicious intent; analysts must corroborate.

# Designed for real operational users



## Guided dashboard

Police officers and field staff

- plain-language priority
- key findings
- recommended next step



## Analyst dashboard

Cyber, intelligence and trust-and-safety teams

- manual and bulk intake
- case drill-down and graph
- score visualizations and live events



## Superuser dashboard

Administrators and technical supervisors

- federation health
- validation and synchronization controls
- system-wide monitoring



**What analysts should see first:** one-sentence assessment, strongest evidence for and against, uncertainty, and the next investigative step.

# Security, provenance and controlled sharing



## Local / sovereign deployment

core text analysis can remain inside agency infrastructure



## Evidence integrity

hashes, fingerprints, signatures, timestamps and source IDs




## Controlled sharing

default redaction, handling tags, signed packages and encrypted payloads



## Graceful degradation

optional models can be unavailable without breaking the workflow



## Interoperability

direction toward STIX 2.1, TAXII 2.1 and TLP 2.0



Hash



Sign



Encrypt



Audit



Share



Provenance supports traceability; it does not by itself prove factual truth.

# Roadmap to validated deployment



**TattvaDrishti** does not replace the analyst —  
*it gives the analyst a **faster, connected** and **more defensible** view of the evidence.*



From SIH-winning prototype to **agency-ready** product.