



TattvaDrishti

Executive Abstract and Project Brief

Prepared for the Director, National Cyber Security Research Council (NCSRC) | July 2026

PROJECT ABSTRACT

TattvaDrishti is an analyst-centric cyber-intelligence and decision-support platform developed by Team ASHTOJ to address the national-security risks created by AI-enabled malign information operations. The platform is designed for police, cyber-security and intelligence teams that must assess large volumes of suspicious digital content without depending on a single opaque detector. It converts raw messages, public-channel content, structured archives and images into traceable analyst cases that preserve source context, evidence lineage and audit history.

The system combines local large-language-model reasoning, transformer-based AI-origin estimation, behavioural manipulation indicators and stylometric evidence to produce an explainable operational priority. It then enriches each case through provenance and fingerprint checks, actor-narrative-region graph analysis, keyword and geographic filtering, and controlled sharing through redacted, signed and encrypted packages. The analytical design is intentionally human-in-the-loop: risk scores prioritize review and surface uncertainty; they do not independently establish truth, identity, intent, guilt or enforcement action.

Scalable evidence intake Manual submissions, JSON/JSONL/NDJSON archives, and normalized Telegram and Reddit content.	Evidence fusion Semantic risk, AI-origin, behavioural and stylometric signals remain visible and are fused into an analyst assessment.
Campaign intelligence Graph analysis links content, actors, narratives, regions and URLs to surface repeated patterns and potential coordination.	Sovereign and auditable operation Core text analysis can run within agency infrastructure; sharing retains hashes, signatures, redaction and audit context.

PROTOTYPE STATUS AND OPERATIONAL VALUE

The current prototype supports role-specific dashboards, real-time case events, threat-intelligence and SIEM-oriented outputs, a geographic risk view and batch processing with per-record error isolation. A recent prototype demonstration on the team's setup processed approximately 150 messages in 10-15 seconds after batching and parallelization improvements. This is an engineering observation, not yet a reproducible benchmark; formal evaluation will document hardware, model configuration, latency percentiles, throughput, failure rates and representative message distributions.

The next stage is a controlled operational pilot rather than additional feature accumulation. Priority work includes authentication and role-based access, hardened storage and job processing, multilingual and code-mixed calibration, scalable graph analytics, STIX/TAXII/TLP-aligned exchange, adversarial testing and structured studies with law-enforcement and intelligence users. Engagement with NCSRC would provide valuable technical review, operational validation and guidance toward a responsible agency-facing pilot.

Project objective
Reduce the time required to reach a defensible analyst assessment while preserving evidence, uncertainty, privacy, auditability and human control.